CSS 534 Assignment 5

En Li, Yangde Li, Yangxiao Wang

December 11, 2018

Contents

1   Application Overview

For project 5, we implemented the Ant Colony Optimization of Traveling Salesman problem. The Ant Colony Optimization imitates real ant colonies. When ants look for food on the ground. a group of ants will each explore different route. If they find food, on their way back, they release Pheromones to mark the route for later searchers. When more ants explore a good routes, Pheromones on the route becomes more intense. Routes with heavier Pheromones will take precedence over routes with lighter Pheromones in the future. To encourage ants to explore different route and avoid local optima, we introduce a evaporation rate to adjust Pheromones at each iteration. We use these information to calculate a probability for each ant to move from city i to city j using the following formula:

$$p_{ij}{}^{k}(t) = \begin{cases} \dfrac{[\tau_{ij}(t)]^{\alpha} \bullet [\eta_{ik}(t)]^{\beta}}{\sum\limits_{s \subset allowed_k} [\tau_{is}(t)]^{\alpha} \bullet [\eta_{is}(t)]^{\beta}}, & \text{if } j \in allowed_k \\ 0 & \text{elsewise} \end{cases} \quad (1)$$

Figure 1: Probability formula

There are three ways to update Pheromones in Ant Colony Optimization: Ant Cycle, Ant Density and Ant Quantity. We used the Ant Cycle model as it is more straight forward. In Ant Cycle model, the change in Pheromone is calculated by diving the current Pheromone by the total distance of the route calculated for current iteration if the ant passed through a given city pair, if the ant did not pass through the city pair, there is no change in Pheromone for this iteration.

2   Parallelization techniques

2.1   MpiJava

In this implementation with MpiJava, we mainly focus on two types of parallelism techniques: **Task** and **Data** decomposition. For the task decomposition, we split the ants to different computing nodes, thus to accelerate whole computing speed. In every optimization loop, the worker nodes send back the delta pheromone generated by the ants in this iteration to master, then master updates the whole pheromone matrix and scatter it to all workers. Another parallel technique is to make every worker run a complete ACO process and update the pheromone matrix using mean value, this is a way to increase the ants searching space and increase the possibility of finding the best path.
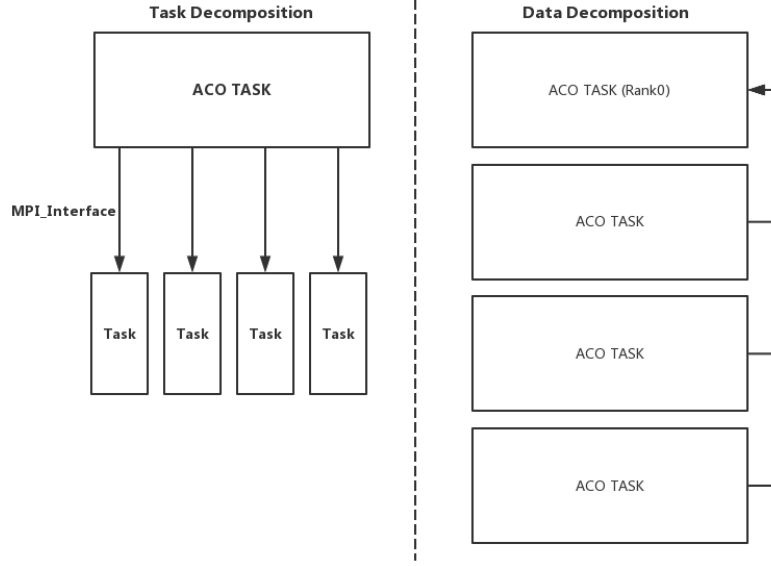
Figure 2: Two types of Parallel Pattern

2.2 MapReduce

At the beginning, I have to state that MapReduce is not the best choice for Traveling Salesman problem with Ant Colony Optimization. One of the most important concept that ACO introduced was pheromones. MapReduce does not support global variables since its data should be distributed in different data nodes. Therefore, maintain and update the pheromones data between each iteration is not feasible with MapReduce or is against the principle of this framework.

There are two strategies to parallelize the ACO algorithm. First one is called chaining MapReduce jobs which contains multiple MapReduce iterations. It stores its intermediate results into binary files and read those files at the beginning of next iteration [1]. The second one takes one-iteration approach. It replicates input files and execute ACO algorithm separately and find the best result among all mappers inside the reducer. The second approach should be better than the first one according to [2] and [3]. The processes that start the MapReduce jobs are time consuming and the number of MapReduce jobs should be as limited as possible. Therefore, we implemented the one-iteration approach with the Ant Colony Optimization algorithm. The number of mapper is determined by the number of input files. Each input file has only one line contains the coordinates of all cities. After reading the file, the first step inside each mapper is parsing the coordinates and calculate the distance between all cities. The calculation result is stored in memory so it can be accessed instantly. Then the ACO process can find the best route for this mapper. After all mapper finished their job, reducer can collect all routes and output the best one.
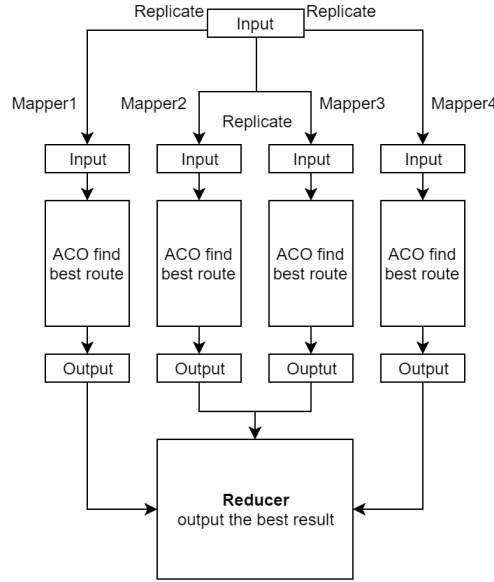
4

Figure 3: MapReduce single iteration

2.3 Spark

In spark version of the project. We took a different approach. In the sequential version of the ant colony optimization. We instantiated a couple 2 dimensional arrays. The natural thing to do is to create a RDD for each matrix. Each of them will be updated through out the program, values in some of the matrices are dependent on other matrices. The approach we took was that we created a new Ant class that implements the Serializable class. Each Ant carries a route array that stores the route for individual ant, a ant ID number to store in ROUTES 2d array, as well as the distance of the route stored in route array. I then created a list of Ant, the size is the number of ants we want to run. I used Spark's parallelize() method to convert the list into a RDD. Now that we have a initial JavaRDD. We can start the optimization. I put all my optimization into a while loop, the loop terminals when the number of interactions we preset has been reached. In the while loop, I first broadcast my Pheromone matrix using Spark broadcast function. This is important because each iteration, the Pheromone is updated using the distance of each ant calculated from previous iteration. But in Spark, we cannot simply store the Pheromone in a global matrix and update them in the RDD. I first transform my ant RDD to a PairRDD where they key is the distance calculated, and the value is a Ant class with updated value. Inside the transformation, I let each ant calculated their own route and calculate the distance. I record the route and distance and return the Tuple. Now that I have a tuple RDD where the key is the distance, I can use RDD sortByKey function to sort the RDD. I obtain the best result from all ants by simply collect the RDD and get the first element. I use a global bestRoute variable to keep track of the best result so far. I then update the pheromone. Note that in updatePheromone function, we need the distance from all Ants to calculate delta Pheromone. We store that into a temporary array called currentLengths. The last transformation is to reset all Tuples in the RDD to original states: set each element in route array for each ant to -1, set their distance to 0, etc.

2.4 MASS

MASS is agent based framework and it perfectly fits the concept of Ant Colony Optimization. Basically, ACO moves ants around the graph to find the best route, and MASS also moves Agents around the Places to process tasks. They have very similar ideas. To parallelize ACO based TSP

with MASS, simply use the graph as Places, and ants as Agents. Just put all of our ACO methods into the callAll method.

```java
@Override
public Object callMethod(int functionId, Object argument)
{
    switch (functionId)
    {
        case INIT: return init();
        case CONNECT: return connectCITIES(argument);
        case SET_POSITION: return setCITYPOSITION(argument);
        case OPTIMIZE: return optimize();
        default:
        return "Unknown Method Number: " + functionId;
    }
}
```

Then, for each iteration,

```java
for (int i = 0; i < ITERATIONS; i++)
{
    // let ants go
    ants.callAll(ACO.OPTIMIZE);
    // update pheronmen
    ants.manageAll();
}
```

3 Execution performance

3.1 MpiJava

For task parallelization, computing task is split into different nodes and the exchanged data is rather small so the cost in communication among nodes is rather small. The performance improvement is obvious. However, the Data Decomposition duplicates the whole data set to run in every nodes, so the performance improvement should be little but the output would be better than sequential version. The 2nd parallel pattern iterate n times data in the same time with the sequential version. Luckily, the Test results approve our speculation.

Figure 4: Execution Snapshots for MpiJava

3.2 MapReduce

The performance of MapReduce is better than we initially thought. The distance of the best route with 1000 iteration is 598. It took 44 seconds. I also tried to increase the number of iterations to 100000, and the result is 569 with 2052 seconds. However, we found the sequential version has better performance than the parallel version. This was caused by the ACO algorithm we implemented. With the similar amount of iterations, the best route it could find was almost the same. When the "mutation rate" for the ACO was increased, the results between different nodes were different, the best route was longer. To improve this, the ACO algorithm has to be adjusted with a different randomize strategy.

Figure 5: Execution output for mapreduce

3.3 Spark

Unfortunately, with 1000 iterations, I only get to 1692 total distances as best route, I let the Spark program run longer and did not obtain any better results. I suspect that this is due to the fact the broadcast function may be slowing the program down by a lot. Every iteration, it needs to be recalculated to rebroadcast to all nodes. In my implementation: I let each RDD element calculate its own route, and I collect the RDD element in every iteration. there are a lot of transforming and collecting over the course of the program, it will slow down the calculation even further.

```
ry on 172.28.203.196:41412 (size: 1831.0 B, free: 366.3 MB)
2018-12-10 17:50:56 INFO  BlockManagerInfo:54 – Added broadcast_0_piece0 in memo
ry on 172.28.203.196:41412 (size: 467.0 B, free: 366.3 MB)
2018-12-10 17:50:56 INFO  BlockManagerInfo:54 – Added broadcast_1_piece0 in memo
ry on 172.28.203.196:41412 (size: 7.2 KB, free: 366.3 MB)
2018-12-10 17:50:56 INFO  TaskSetManager:54 – Starting task 1.0 in stage 0.0 (TI
D 1, 172.28.203.196, executor 0, partition 1, PROCESS_LOCAL, 9467 bytes)
2018-12-10 17:50:56 INFO  TaskSetManager:54 – Finished task 0.0 in stage 0.0 (TI
D 0) in 623 ms on 172.28.203.196 (executor 0) (1/2)
2018-12-10 17:50:56 INFO  TaskSetManager:54 – Finished task 1.0 in stage 0.0 (TI
D 1) in 56 ms on 172.28.203.196 (executor 0) (2/2)
2018-12-10 17:50:56 INFO  TaskSchedulerImpl:54 – Removed TaskSet 0.0, whose task
s have all completed, from pool
2018-12-10 17:50:56 INFO  DAGScheduler:54 – ResultStage 0 (collect at ACOSpark.j
ava:282) finished in 2.068 s
2018-12-10 17:50:56 INFO  DAGScheduler:54 – Job 0 finished: collect at ACOSpark.
java:282, took 2.154926 s
distance is 1691.9204807203478
2018-12-10 17:50:56 INFO  SparkContext:54 – Starting job: collect at ACOSpark.ja
va:283
2018-12-10 17:50:56 INFO  DAGScheduler:54 – Got job 1 (collect at ACOSpark.java:
283) with 2 output partitions
```

Figure 6: Execution output for spark

3.4 MASS

The parallel strategy of multi-agents is similar with ACO. MASS is based on multi-thread, and we mainly used Agent to present ants. The number of ants is small and most of the thread overhead comes from the number of iteration. As a result, the performance is nearly the same as the sequential version.

Figure 7: Execution output for MASS

3.5 Overall Performance

The following table compares results from each parallel technique with the same parameters and 1000 iterations. The spark's performance results are not included the graph, because there are significant differences between spark's results and the others.

Figure 8: Performance Comparison with 4 ants

And we could not get stable results because other teams were using the same remote computer, however, the graph still shows that the elapsed time increased with the number of iterations. We can also observe that the performance of MapReduce is the highest compare to the other 3.

9

Iterations	Sequential	MpiJava	MapReduce	Spark	MASS
100	1.504	2.030	9.047	109	7.344
200	3.740	2.637	30.397	210	18.904
300	6.715	5.330	26.909	355	10.288
400	8.517	4.673	84.700	409	26.329
500	9.461	6.946	97.053	510	17.706
600	18.144	10.716	127.428	599	18.961
700	32.685	10.275	102.772	830	41.059
800	40.459	17.673	117.578	989	21.622
900	38.326	21.474	58.927	1020	33.032
1000	45.287	17.019	66.191	1140	47.465
Best Route	638.1	911.37	602.93	1692	777.65

Table 1: Performance Comparison with 4 ants

We took 4 spatial scalability and used MpiJava. From the graph, we can observe that the ACO's scalability is not ideal. This is caused by the algorithm we used. With more computation nodes, the ants are still runs in the same graph. We also observed the same results on other platform.
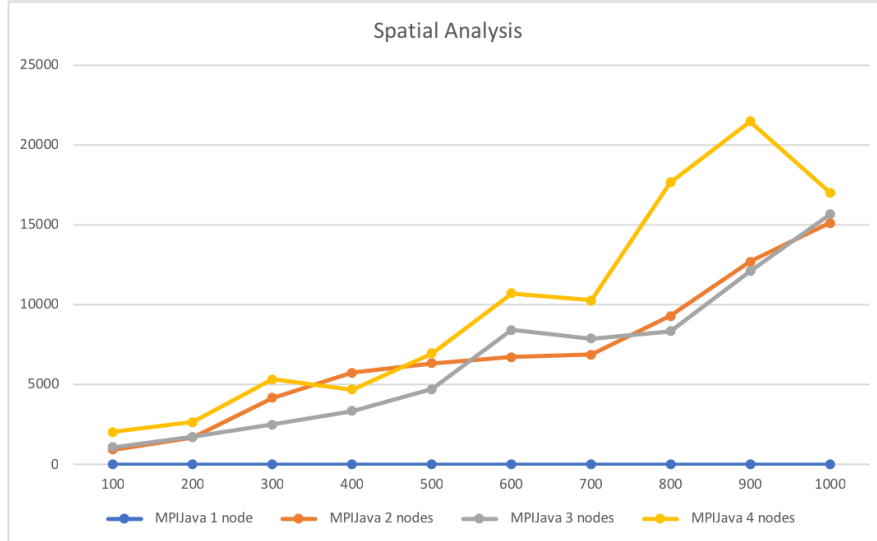
Figure 9: Spatial analysis

4 Programmability analysis

4.1 MpiJava

As for the programmability, MPI keeps a low invasion in original code. The MPI context code is only several numbers of line. However, the MPI API is relatively complicated. The communication API like send/recv contains a long parameter list and this kind of programming style is unfriendly to new developers. The users have to look up the api docs for the function calls they want to invoke, and the whole program is split by many rank branches which leads to a unreadable code base.

10

4.2   MapReduce

With MapRedour, the implementation of ACO is straightforward, every mapper runs its own ACO and report a best result at the end of all iteration. Thanks to the framework, only Mapper and Reducer class need to be implemented, the amount of boilerplate is minimal compare to other method.

4.3   Spark

Our specific implementation of the ACO may not be the best fit to be run on Spark. In general, one main advantage of Spark is the amount of boiler code can be fairly small due to the advantage of using Lamda expression. But in this program, I actually had to write more lines of code than the sequential version. The extra boilerplate code comes mainly from the work to rebroadcast Pheromone. I wrote code to broadcast it, collect distances from all ants, and then rebroadcast it after updating the Pheromone. This is because in Spark, we cannot update global data structures inside of RDD. We can read them after broadcasting it. To change a shared data structure across different nodes. I had to collect data needed to update the Pheromone in every iteration before updating it, and then rebroadcast.

4.4   MASS

There is a **ACO** class in our implementation of the ACO. This class contains all ACO algorithm such as updating pheromone or move ants. To convert it to MASS, simply extend it to **Agent** and override *callMethod* function. And make some adjustments to fit the MASS api. Therefore, there is not a large amount of boilerplate code and minor changes compare to our base sequential version.

4.5   Overall Programmability

-	Sequential	MpiJava	MapReduce	Spark	MASS
Proximate boilerplate ratio	0	0.2	0.2	0.5	0.3

Table 2: Proximate boilerplate ratio

5   Lab 5


```
[wyxiao_css534@cssmpi1 prog5]$ java -jar prog5-1.0-SNAPSHOT.jar
MASS.init: done
Execution time = 770 milliseconds
[wyxiao_css534@cssmpi1 prog5]$
```

Figure 10: Execution output for lab5 run sample

```java
package edu.uwb.css534;

import edu.uw.bothell.css.dsl.MASS.*; // Library for Multi-Agent Spatial Simulation
import edu.uw.bothell.css.dsl.MASS.logging.LogLevel;
import edu.uw.bothell.css.dsl.MASS.Agents;
import edu.uw.bothell.css.dsl.MASS.MASS;
import edu.uw.bothell.css.dsl.MASS.Places;
import java.util.*;
import java.io.*;
import java.lang.reflect.Array;
import java.util.Arrays;

public class lab5
{

    private static final String NODE_FILE = "nodes.xml";


    public static void main(String[] args) throws Exception
    {


        MASS.setNodeFilePath(NODE_FILE);
        MASS.setLoggingLevel(LogLevel.DEBUG);
        MASS.init();

        Places matrix = new Places( 1, Matrix.class.getName(), null, 10, 10);
        Agents worker = new Agents( 2, Worker.class.getName(), null, matrix, 2);

        worker.callAll( Worker.goElsewhere_ );
        worker.manageAll( );

        MASS.finish();
    }

}


public class Matrix extends Place
{
    public Matrix()
    {

    }
    public Matrix(Object obj)
    {

    }
    public Object callMethod(int method, Object o)
    {
        switch (method)
        {
            default:
            return new String("Unknown Method Number: " + method);
        }
    }
}

public class Worker extends Agent
{

    public static final int goElsewhere_ = 0;

    public Worker()
    {
        super();
    }
    public Worker(Object object)
    {
```

```
70          super();
71      }
72      public Object callMethod(int funcId)
73      {
74          switch (funcId)
75          {
76              case goElsewhere_:
77              return goElsewhere();
78          }
79          return null;
80      }
81      public Object goElsewhere()
82      {
83          int newX = 0; // a new destination's X-coordinate
84          int newY = 0; // a new destination's Y-coordinate
85          int min = 1; // a new destination's # agents
86
87          int currX = getPlace().getIndex()[0], currY = getPlace().getIndex()[1];
88          int sizeX = getPlace().getSize()[0], sizeY = getPlace().getSize()[1];
89
90          Random generator = new Random();
91          boolean candidatePicked = false;
92
93          int next = 0;
94          next = generator.nextInt(1);
95          if (next == 1)
96          {
97              newX = currX + generator.nextInt(sizeX - currX - 1);
98          }
99          else
100         {
101             newX = currX - generator.nextInt(currX);
102         }
103
104         next = generator.nextInt(1);
105         if (next == 1)
106         {
107             newY = currY + generator.nextInt(sizeY - sizeY - 1);
108         }
109         else
110         {
111             newY = currY - generator.nextInt(currY);
112         }
113
114         migrate(newX, newY);
115         return null;
116     }
117 }
```

References

[1] A. Mohan and R. G, "A Parallel Implementation of Ant Colony Optimization for TSP based on MapReduce Framework," *International Journal of Computer Applications*, vol. 88, no. 8, pp. 9–12, 2014.

[2] B. Wu, G. Wu, and M. Yang, "A MapReduce based Ant Colony Optimization approach to combinatorial optimization problems," *2012 8th International Conference on Natural Computation*, 2012.

[3] G. Cesari. Divi de and conquer strategies for parallel tsp heuristics. Comput. Oper. Res., 23:681–694, July 1996