

Program 1 Report

Yangxiao Wang

Contents

1	Documentation	2
2	Source code	2
2.1	Wave2D.cpp	2
2.2	Wave2D mpi.cpp	5
3	Execution output	5
3.1	Output analysis	5
3.2	Execution output	5
4	Discussions	5
5	Lab Sessions 2	5
5.1	Source Code	6
5.2	Execution output	8

1 Documentation

In this project, we tried to solve Traveling salesman problem (TSP) with genetic algorithm. We are required to implement the core part of the algorithm: evaluation, crossover and mutation. Tsp problem often takes large amount of data and time. And to improve the efficiency, our approach is using OpenMP. My initial attempt is implement a greedy parallelization: parallelize the whole program from *evaluate()* to *populate()* with:

```
1  #pragma omp parallel for
```

however, after trying different things with the program, I realized that populate and mutate does not need to be parallelized. The time spent did increase not have much difference. And *select()* with parallelization could spend more time than when running multi-thread. Therefore, the only functions I need to parallelize are *crossover()* and *evaluate()* which contain much larger loops. I simply added `#pragma omp parallel for` before the outer for loop inside *crossover()* and *evaluate()*.

2 Source code

2.1 Wave2D.cpp

```
1  #include <iostream>
2  #include "Timer.h"
3  #include <stdlib.h>    // atoi
4  #include <math.h>
5  #include <stdio.h>
6
7  int default_size = 100; // the default system size
8  int defaultCellWidth = 8;
9  double c = 1.0;        // wave speed
10 double dt = 0.1;        // time quantum
11 double dd = 2.0;        // change in system
12
13 using namespace std;
14
15 int main(int argc, char *argv[]) {
16     // verify arguments
17     if (argc != 4) {
18         cerr << "usage: Wave2D size max.time interval" << endl;
19         return -1;
20     }
21     int size = atoi(argv[1]);
22     int max_time = atoi(argv[2]);
23     int interval = atoi(argv[3]);
24
25     if (size < 100 || max_time < 3 || interval < 0) {
26         cerr << "usage: Wave2D size max.time interval" << endl;
27         cerr << "      where size >= 100 && time >= 3 && interval >= 0" << endl;
28         return -1;
29     }
30
31     // create a simulation space
32     double z[3][size][size];
33     for (int p = 0; p < 3; p++)
34         for (int i = 0; i < size; i++)
35             for (int j = 0; j < size; j++)
```

```

36         z[p][i][j] = 0.0; // no wave
37
38     // start a timer
39     Timer time;
40     time.start();
41
42     // time = 0;
43     // initialize the simulation space: calculate z[0][][]
44     int weight = size / default_size;
45     for (int i = 0; i < size; i++) {
46         for (int j = 0; j < size; j++) {
47             if (i > 40 * weight && i < 60 * weight &&
48                 j > 40 * weight && j < 60 * weight) {
49                 z[0][i][j] = 20.0;
50             } else {
51                 z[0][i][j] = 0.0;
52             }
53         }
54     }
55
56     // time = 1
57     for (int i = 1; i < size - 1; i++) {
58         for (int j = 1; j < size - 1; j++) {
59             z[1][i][j] = z[0][i][j] + (pow(c, 2) / 2) * pow(dt / dd, 2) * (z[0][i + 1][j] + z[0][i - 1][j] + z[0][i][j + 1] + z[0][i][j - 1] - (4.0 * z[0][i][j]));
60         }
61     }
62
63     // simulate wave diffusion from time = 2
64     for (int t = 2; t < max_time; t++) {
65         int time = t % 3;
66         for (int i = 1; i < size - 1; i++) {
67             for (int j = 1; j < size - 1; j++) {
68                 int time_1, time_2;
69                 //rotate z
70                 if (time == 0) {
71                     time_1 = 2;
72                     time_2 = 1;
73                 } else if (time == 1) {
74                     time_1 = 0;
75                     time_2 = 2;
76                 } else {
77                     time_1 = 1;
78                     time_2 = 0;
79                 }
80                 //calculation
81                 z[time][i][j] =
82                 2.0 * z[time_1][i][j] - z[time_2][i][j] + (pow(c, 2) * pow(dt / dd, 2) * (z[time_1][i + 1][j] + z[time_1][i - 1][j] + z[time_1][i][j + 1] + z[time_1][i][j - 1] - (4.0 * z[time_1][i][j])));
83             }
84         }
85         //print out
86         if (interval != 0 && t % interval == 0) {
87             cout << t << endl;
88             for (int j = 0; j < size; j++) {
89                 for (int i = 0; i < size; i++) {
90                     cout << z[time][i][j] << " ";
91                 }
92                 cout << endl;

```

```
93         }
94         cout << endl;
95     }
96 } // end of simulation
97
98 // finish the timer
99 cerr << "Elapsed time = " << time.lap() << endl;
100 return 0;
101 }
```

2.2 Wave2D mpi.cpp

3 Execution output

3.1 Output analysis

1. The shortest trip in my program is equal to 447.638
2. The performance improvement with four threads in my program is equal to $50548672 / 23005048 = 2.2$ times

3.2 Execution output

Check if output is correct

```
1 [wyxiao_css534@cssmpi1 prog2]$ ./Wave2D 576 500 50 > reS.txt
2 [wyxiao_css534@cssmpi1 prog2]$ mpirun -n 4 ./Wave2D_mpi 576 500 50 1 > reF.txt
3 [wyxiao_css534@cssmpi1 prog2]$ diff reF.txt reS.txt
4 [wyxiao_css534@cssmpi1 prog2]$
```

Check output the performance improvement with four machines: $5721050 / 1575955 = 3.63$ times

```
1 [wyxiao_css534@cssmpi1 prog2]$ ./Wave2D 576 500 0
2 Elapsed time = 5721050
3 [wyxiao_css534@cssmpi1 prog2]$ mpirun -n 4 ./Wave2D_mpi 576 500 0 1
4 Elapsed time = 1575955
```

4 Discussions

In addition, I tried to improve the whole efficiency by replacing calculating distance with a cached matrix that contains all the $36 * 36$ distance. This implementation decrease significant amount of time spent, although this made the program program with single thread runs faster than that with multi-thread. This also explained why select() with multi-thread could be slower than that with single thread. Starting a multi-thread could require some time to analysis the for-loop and distribute the tasks. Therefore, it is best to use multi-thread when the loop is large and require large computational power.

To improve the performance of current program, just use the implementation I mentioned above to shorten the time required to calculate the distance. However, this only works with the current data set, it could be less efficient with larger data sets.

5 Lab Sessions 2

Lab 2 we parallelize two programs that compute Pi using Monte Carlo methods and Integration. As the result shown, multi-thread decrease significant amount of time that required to finish the program. However, if the number of iterations is too small, the performance would decrease.

5.1 Source Code

```
1  #include "mpi.h"
2  #include <stdlib.h> // atoi
3  #include <iostream> // cerr
4  #include "Timer.h"
5
6  using namespace std;
7
8  void init(double *matrix, int size, char op) {
9      for (int i = 0; i < size; i++)
10         for (int j = 0; j < size; j++)
11             matrix[i * size + j] = (op == '+') ? i + j : ((op == '-') ? i - j : 0);
12 }
13
14 void print(double *matrix, int size, char id) {
15     for (int i = 0; i < size; i++)
16         for (int j = 0; j < size; j++)
17             cout << id << "[" << i << "]"[" << j << "]" = " << matrix[i * size + j] << endl;
18 }
19
20 void multiplication(double *a, double *b, double *c, int stripe, int size) {
21     for (int k = 0; k < size; k++)
22         for (int i = 0; i < stripe; i++)
23             for (int j = 0; j < size; j++)
24                 // c[i][k] += a[i][j] * b[j][k];
25                 c[i * size + k] += a[i * size + j] * b[j * size + k];
26 }
27
28 int main(int argc, char *argv[]) {
29     int my_rank = 0; // used by MPI
30     int mpi_size = 1; // used by MPI
31     int size = 400; // array size
32     bool print_option = false; // print out c[] if it is true
33     Timer timer;
34
35     // variables verification
36     if (argc == 3) {
37         if (argv[2][0] == 'y')
38             print_option = true;
39     }
40
41     if (argc == 2 || argc == 3) {
42         size = atoi(argv[1]);
43     } else {
44         cerr << "usage:    matrix size [y|n]" << endl;
45         cerr << "example: matrix 400    y" << endl;
46         return -1;
47     }
48
49     MPI_Init(&argc, &argv); // start MPI
50     MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
51     MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);
52
53     // matrix initialization
54     double *a = new double[size * size];
55     double *b = new double[size * size];
56     double *c = new double[size * size];
57 }
```

```

58     if (my_rank == 0) { // master initializes all matrices
59         init(a, size, '+');
60         init(b, size, '-');
61         init(c, size, '0');
62
63         // print initial values
64         if (false) {
65             print(a, size, 'a');
66             print(b, size, 'b');
67         }
68
69         // start a timer
70         timer.start();
71     } else { // slaves zero-initializes all matrices
72         init(a, size, '0');
73         init(b, size, '0');
74         init(c, size, '0');
75     }
76
77     // broadcast the matrix size to all.
78     MPI_Bcast(&size, 1, MPI_INT, 0, MPI_COMM_WORLD);
79
80     int stripe = size / mpi_size; // partitioned stripe
81
82     // master sends each partition of a[] to a different slave
83     // master also sends b[] to all slaves
84     if (my_rank == 0) {
85         for (int rank = 1; rank < mpi_size; ++rank) {
86             MPI_Send(a + rank * stripe * size, size * stripe, MPI_DOUBLE, rank, 0, ←
87                     MPI_COMM_WORLD);
88             MPI_Send(b, size * size, MPI_DOUBLE, rank, 0, MPI_COMM_WORLD);
89         }
90     } else {
91         MPI_Status status;
92         MPI_Recv(a, size * stripe, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
93         MPI_Recv(b, size * size, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
94     }
95     multiplication(a, b, c, stripe, size); // all ranks should compute ←
96     // multiplication
97
98     // master receives each partition of c[] from a different slave
99     if (my_rank == 0) {
100         for (int rank = 1; rank < mpi_size; ++rank) {
101             MPI_Status status;
102             MPI_Recv(c + rank * stripe * size, size * stripe, MPI_DOUBLE, rank, 0, ←
103                     MPI_COMM_WORLD,
104                     &status);
105         }
106     } else {
107         MPI_Send(c, stripe * size, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
108     }
109
110     if (my_rank == 0)
111         // stop the timer
112         cout << "elapsed time = " << timer.lap() << endl;
113
114     // results
115     if (print_option && my_rank == 0)
116         print(c, size, 'c');

```

```
116     MPI_Finalize(); // shut down MPI
117 }
```

5.2 Execution output

Check output is correct

```
1 [wyxiao_css534@cssmpi1 lab2]$ ./matrix 100 y > maxS.txt
2 [wyxiao_css534@cssmpi1 lab2]$ mpirun -n 4 ./matrix_mpi 100 y > maxM.txt
3 [wyxiao_css534@cssmpi1 lab2]$ diff maxS.txt maxM.txt
4 lc1
5 < elapsed time = 5185
6 ____
7 > elapsed time = 7133
8 [wyxiao_css534@cssmpi1 lab2]$
```

Check output the performance improvement: $774693 / 258639 = 2.9952$ times

```
1 [wyxiao_css534@cssmpi1 lab2]$ ./matrix 500
2 elapsed time = 774693
3 [wyxiao_css534@cssmpi1 lab2]$ mpirun -n 4 ./matrix_mpi 500
4 elapsed time = 258639
5 [wyxiao_css534@cssmpi1 lab2]$
```