

Program 2 Report

Yangxiao Wang

Contents

1	Documentation	2
2	Source code	2
2.1	Wave2D.cpp	2
2.2	Wave2D mpi.cpp	5
3	Execution output	8
3.1	Output analysis	8
3.2	Execution output	8
4	Discussions	8
5	Lab Sessions 2	8
5.1	Source Code	9
5.2	Execution output	11

1 Documentation

In this project, we tried to build a two-dimensional wave diffusion program by using Schroedinger's Wave Dissemination formula. Since the formula requires heavy computational power, we proposed to use parallelization techniques to implement it. And in this report, I would demonstrate the code and time required for each of the following implementations: sequential, MPI, and hybrid form of MPI and OpenMPI.

The MPI methods I used is the two blocking communications *MPI_Send()* and *MPI_Recv()*. The advantage of using blocking communication is I do not need to handle the synchronization problem. The program can only continue to work until all communication is done (all data is sent and received). In addition, to exchange the boundary data, my approach was sending the boundary stripe at beginning of every loop of t (time). And for implementation of OpenMPI, I simply parallelized the Schroedinger's formula part. For each ranks, it has multiple threads to compute the formula.

2 Source code

2.1 Wave2D.cpp

```
1  #include <iostream>
2  #include "Timer.h"
3  #include <stdlib.h>    // atoi
4  #include <math.h>
5  #include <stdio.h>
6
7  int default_size = 100; // the default system size
8  int defaultCellWidth = 8;
9  double c = 1.0;        // wave speed
10 double dt = 0.1;        // time quantum
11 double dd = 2.0;        // change in system
12
13 using namespace std;
14
15 int main(int argc, char *argv[]) {
16     // verify arguments
17     if (argc != 4) {
18         cerr << "usage: Wave2D size max_time interval" << endl;
19         return -1;
20     }
21     int size = atoi(argv[1]);
22     int max_time = atoi(argv[2]);
23     int interval = atoi(argv[3]);
24
25     if (size < 100 || max_time < 3 || interval < 0) {
26         cerr << "usage: Wave2D size max_time interval" << endl;
27         cerr << "      where size >= 100 && time >= 3 && interval >= 0" << endl;
28         return -1;
29     }
30
31     // create a simulation space
32     double z[3][size][size];
33     for (int p = 0; p < 3; p++)
34         for (int i = 0; i < size; i++)
35             for (int j = 0; j < size; j++)
36                 z[p][i][j] = 0.0; // no wave
```

```

37
38 // start a timer
39 Timer time;
40 time.start();
41
42 // time = 0;
43 // initialize the simulation space: calculate z[0][][]
44 int weight = size / default_size;
45 for (int i = 0; i < size; i++) {
46     for (int j = 0; j < size; j++) {
47         if (i > 40 * weight && i < 60 * weight &&
48             j > 40 * weight && j < 60 * weight) {
49             z[0][i][j] = 20.0;
50         } else {
51             z[0][i][j] = 0.0;
52         }
53     }
54 }
55
56 // time = 1
57 for (int i = 1; i < size - 1; i++) {
58     for (int j = 1; j < size - 1; j++) {
59         z[1][i][j] = z[0][i][j] + (pow(c, 2) / 2) * pow(dt / dd, 2) * (z[0][i + ←
60             1][j] + z[0][i - 1][j] + z[0][i][j + 1] + z[0][i][j - 1] - (4.0 * ←
61             z[0][i][j]));
62     }
63 }
64
65 // simulate wave diffusion from time = 2
66 for (int t = 2; t < max_time; t++) {
67     int time = t % 3;
68     for (int i = 1; i < size - 1; i++) {
69         for (int j = 1; j < size - 1; j++) {
70             int time_1, time_2;
71             //rotate z
72             if (time == 0) {
73                 time_1 = 2;
74                 time_2 = 1;
75             } else if (time == 1) {
76                 time_1 = 0;
77                 time_2 = 2;
78             } else {
79                 time_1 = 1;
80                 time_2 = 0;
81             }
82             //calculation
83             z[time][i][j] =
84             2.0 * z[time_1][i][j] - z[time_2][i][j] + (pow(c, 2) * pow(dt / dd, ←
85                 2) * (z[time_1][i + 1][j] + z[time_1][i - 1][j] + z[time_1][i][←
86                 j + 1] + z[time_1][i][j - 1] - (4.0 * z[time_1][i][j])));
87         }
88     }
89     //print out
90     if (interval != 0 && t % interval == 0) {
91         cout << t << endl;
92         for (int j = 0; j < size; j++) {
93             for (int i = 0; i < size; i++) {
94                 cout << z[time][i][j] << " ";
95             }
96             cout << endl;
97         }
98     }
99 }

```

```
94         cout << endl;
95     }
96 } // end of simulation
97
98 // finish the timer
99 cerr << "Elapsed time = " << time.lap() << endl;
100 return 0;
101 }
```

2.2 Wave2D mpi.cpp

```
1  #include <iostream>
2  #include "Timer.h"
3  #include <stdlib.h>    // atoi
4  #include <math.h>
5  #include <mpi.h>
6  #include <stdio.h>
7  #include <omp.h>
8
9  int default_size = 100; // the default system size
10 int defaultCellWidth = 8;
11 double c = 1.0;        // wave speed
12 double dt = 0.1;       // time quantum
13 double dd = 2.0;       // change in system
14
15 using namespace std;
16
17 int main(int argc, char *argv[]) {
18     int my_rank = 0;      // used by MPI
19     // used by MPI
20
21     // verify arguments
22     if (argc != 5) {
23         cerr << "usage: Wave2D size max_time interval n_thread" << endl;
24         return -1;
25     }
26     int size = atoi(argv[1]);
27     int max_time = atoi(argv[2]);
28     int interval = atoi(argv[3]);
29     int nThreads = atoi(argv[4]);
30     int mpi_size;
31
32     if (size < 100 || max_time < 3 || interval < 0 || nThreads <= 0) {
33         cerr << "usage: Wave2D size max_time interval" << endl;
34         cerr << "where size >= 100 && time >= 3 && interval >= 0 && mpi-size > 0" << endl;
35         return -1;
36     }
37
38     // start MPI
39     MPI_Init(&argc, &argv);
40     MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
41     MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);
42     // change # of threads for openMP
43     omp_set_num_threads(nThreads);
44
45     // create a simulation space
46     double z[3][size][size];
47     for (int p = 0; p < 3; p++)
48         for (int i = 0; i < size; i++)
49             for (int j = 0; j < size; j++)
50                 z[p][i][j] = 0.0; // no wave
51
52     // start a timer
53     Timer time;
54     time.start();
55
56     // time = 0;
57     // initialize the simulation space: calculate z[0][][]
```

```

58     int weight = size / default_size;
59     for (int i = 0; i < size; i++) {
60         for (int j = 0; j < size; j++) {
61             if (i > 40 * weight && i < 60 * weight && j > 40 * weight && j < 60 * weight) {
62                 z[0][i][j] = 20.0;
63             } else {
64                 z[0][i][j] = 0.0;
65             }
66         }
67     }
68
69     // time = 1, and parallelization
70     #pragma omp parallel for
71     for (int i = 1; i < size - 1; i++) {
72         for (int j = 1; j < size - 1; j++) {
73             z[1][i][j] = z[0][i][j] + (pow(c, 2) / 2) * pow(dt / dd, 2) * (z[0][i + 1][j] + z[0][i - 1][j] + z[0][i][j + 1] + z[0][i][j - 1] - (4.0 * z[0][i][j]));
74         }
75     }
76
77
78     int stripe = size / mpi_size; // partitioned stripe
79
80     // simulate wave diffusion from time = 2
81     for (int t = 2; t < max_time; t++) {
82         //rotate z
83         int time = t % 3;
84         int time_1, time_2;
85         if (time == 0) {
86             time_1 = 2;
87             time_2 = 1;
88         } else if (time == 1) {
89             time_1 = 0;
90             time_2 = 2;
91         } else {
92             time_1 = 1;
93             time_2 = 0;
94         }
95
96         //exchange boundary data
97         if (my_rank == 0) {
98             MPI_Send(*(z + time_1) + stripe * (my_rank + 1) - 1), size, MPI_DOUBLE, my_rank + 1, 0, MPI_COMM_WORLD);
99
100             MPI_Status status;
101             MPI_Recv(*(z + time_1) + stripe), size, MPI_DOUBLE, my_rank + 1, 0, MPI_COMM_WORLD, &status);
102         } else if (my_rank == mpi_size - 1) {
103             MPI_Send(*(z + time_1) + stripe * my_rank), size, MPI_DOUBLE, my_rank - 1, 0, MPI_COMM_WORLD);
104
105             MPI_Status status;
106             MPI_Recv(*(z + time_1) + stripe * my_rank - 1), size, MPI_DOUBLE, my_rank - 1, 0, MPI_COMM_WORLD, &status);
107         } else {
108             MPI_Send(*(z + time_1) + stripe * my_rank), size, MPI_DOUBLE, my_rank - 1, 0, MPI_COMM_WORLD);
109             MPI_Send(*(z + time_1) + stripe * (my_rank + 1) - 1), size, MPI_DOUBLE, my_rank + 1, 0, MPI_COMM_WORLD);

```

```

110
111     MPI_Status status;
112     MPI_Recv(*(z + time_1) + stripe * my_rank - 1), size, MPI_DOUBLE, ←
        my_rank - 1, 0, MPI_COMM_WORLD, &status);
113     MPI_Recv(*(z + time_1) + stripe * (my_rank + 1)), size, MPI_DOUBLE, ←
        my_rank + 1, 0, MPI_COMM_WORLD, &status);
114 }
115
116 //Parallelization for the Schroedinger's formula
117 #pragma omp parallel for
118 for (int i = my_rank * stripe; i < (my_rank + 1) * stripe; i++) {
119     if (i == 0 || i == size - 1) {
120         continue;
121     }
122     for (int j = 1; j < size - 1; j++) {
123
124         z[time][i][j] =
125             2.0 * z[time_1][i][j] - z[time_2][i][j] + (pow(c, 2) * pow(dt / dd, ←
                2)
126             * (z[time_1][i + 1][j] + z[time_1][i - 1][j] + z[time_1][i][j + 1] +
127             z[time_1][i][j - 1] - (4.0 * z[time_1][i][j])));
128     }
129 }
130
131 //output if it's interval
132 if (interval != 0 && t % interval == 0) {
133     //Aggregate all results from all ranks
134     if (my_rank == 0) {
135         for (int rank = 1; rank < mpi_size; ++rank) {
136             MPI_Status status;
137             MPI_Recv(*(z + time) + rank * stripe), stripe * size, ←
                MPI_DOUBLE, rank, 0, MPI_COMM_WORLD, &status);
138         }
139
140         cout << t << endl;
141         for (int j = 0; j < size; j++) {
142             for (int i = 0; i < size; i++) {
143                 cout << z[time][i][j] << " ";
144             }
145             cout << endl;
146         }
147         cout << endl;
148
149     } else {
150         MPI_Send(*(z + time) + my_rank * stripe), stripe * size, ←
            MPI_DOUBLE, 0, 0,
151         MPI_COMM_WORLD);
152     }
153 }
154 } // end of simulation
155
156 MPI_Finalize(); // shut down MPI
157
158 // finish the timer
159 if(my_rank == 0) {
160     cerr << "Elapsed time = " << time.lap() << endl;
161 }
162
163 return 0;
164 }

```

3 Execution output

3.1 Output analysis

1. The performance improvement with four machines: $5721050 / 1575955 = 3.63$ times
2. The performance improvement with four machines with multithreading: $5721050 / 874590 = 6.54$ times

3.2 Execution output

Check if output is correct

```
1 [wyxiao_css534@cssmpi1 prog2]$ ./Wave2D 576 500 50 > reS.txt
2 Elapsed time = 7313114
3 [wyxiao_css534@cssmpi1 prog2]$ mpirun -n 4 ./Wave2D_mpi 576 500 50 4 > reF.txt
4 Elapsed time = 6991195
5 [wyxiao_css534@cssmpi1 prog2]$ diff reF.txt reS.txt
6 [wyxiao_css534@cssmpi1 prog2]$
```

Check output the performance improvement with four machines: $5721050 / 1575955 = 3.63$ times

```
1 [wyxiao_css534@cssmpi1 prog2]$ ./Wave2D 576 500 0
2 Elapsed time = 5721050
3 [wyxiao_css534@cssmpi1 prog2]$ mpirun -n 4 ./Wave2D_mpi 576 500 0 1
4 Elapsed time = 1575955
```

Check output the performance improvement with four machines with multithreading: $5721050 / 874590 = 6.54$ times

```
1 [wyxiao_css534@cssmpi1 prog2]$ mpirun -n 4 ./Wave2D_mpi 576 500 0 4
2 Elapsed time = 874590
```

4 Discussions

I noticed the cost of communication is not ideal, especially the aggregation of all ranks' data. Therefore I should keep the numbers of aggregation as low as possible, the aggregation only happens when the output is required.

To improve the current implementation, the best way I could think of is using unblocking communication (*MPI_Isend* and *MPI_Irecv*) when aggregation happens. Then the program would not stop when perform aggregation. And add a *MPI_Wait* statement before $t = t_{current} + 3$ to make sure previous communication is finished when changing the values. Since the aggregation is the most time consuming part.

5 Lab Sessions 2

Lab 2 we parallelize a programs that compute the result of multiplication of matrix using sequential and MPI. As the result shown, four MPI ranks decrease significant amount of time that required to finish the program. However, if the size of matrix is too small, the performance would decrease as shown in the first execution output.

5.1 Source Code

```
1  #include "mpi.h"
2  #include <stdlib.h> // atoi
3  #include <iostream> // cerr
4  #include "Timer.h"
5
6  using namespace std;
7
8  void init(double *matrix, int size, char op) {
9      for (int i = 0; i < size; i++)
10         for (int j = 0; j < size; j++)
11             matrix[i * size + j] = (op == '+') ? i + j : ((op == '-') ? i - j : 0);
12 }
13
14 void print(double *matrix, int size, char id) {
15     for (int i = 0; i < size; i++)
16         for (int j = 0; j < size; j++)
17             cout << id << "[" << i << "]"[" << j << "]" = " << matrix[i * size + j] << endl;
18 }
19
20 void multiplication(double *a, double *b, double *c, int stripe, int size) {
21     for (int k = 0; k < size; k++)
22         for (int i = 0; i < stripe; i++)
23             for (int j = 0; j < size; j++)
24                 // c[i][k] += a[i][j] * b[j][k];
25                 c[i * size + k] += a[i * size + j] * b[j * size + k];
26 }
27
28 int main(int argc, char *argv[]) {
29     int my_rank = 0; // used by MPI
30     int mpi_size = 1; // used by MPI
31     int size = 400; // array size
32     bool print_option = false; // print out c[] if it is true
33     Timer timer;
34
35     // variables verification
36     if (argc == 3) {
37         if (argv[2][0] == 'y')
38             print_option = true;
39     }
40
41     if (argc == 2 || argc == 3) {
42         size = atoi(argv[1]);
43     } else {
44         cerr << "usage:    matrix size [y|n]" << endl;
45         cerr << "example: matrix 400 y" << endl;
46         return -1;
47     }
48
49     MPI_Init(&argc, &argv); // start MPI
50     MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
51     MPI_Comm_size(MPI_COMM_WORLD, &mpi_size);
52
53     // matrix initialization
54     double *a = new double[size * size];
55     double *b = new double[size * size];
56     double *c = new double[size * size];
57 }
```

```

58     if (my_rank == 0) { // master initializes all matrices
59         init(a, size, '+');
60         init(b, size, '-');
61         init(c, size, '0');
62
63         // print initial values
64         if (false) {
65             print(a, size, 'a');
66             print(b, size, 'b');
67         }
68
69         // start a timer
70         timer.start();
71     } else { // slaves zero-initializes all matrices
72         init(a, size, '0');
73         init(b, size, '0');
74         init(c, size, '0');
75     }
76
77     // broadcast the matrix size to all.
78     MPI_Bcast(&size, 1, MPI_INT, 0, MPI_COMM_WORLD);
79
80     int stripe = size / mpi_size; // partitioned stripe
81
82     // master sends each partition of a[] to a different slave
83     // master also sends b[] to all slaves
84     if (my_rank == 0) {
85         for (int rank = 1; rank < mpi_size; ++rank) {
86             MPI_Send(a + rank * stripe * size, size * stripe, MPI_DOUBLE, rank, 0, ←
87                     MPI_COMM_WORLD);
88             MPI_Send(b, size * size, MPI_DOUBLE, rank, 0, MPI_COMM_WORLD);
89         }
90     } else {
91         MPI_Status status;
92         MPI_Recv(a, size * stripe, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
93         MPI_Recv(b, size * size, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &status);
94     }
95     multiplication(a, b, c, stripe, size); // all ranks should compute ←
96     // multiplication
97
98     // master receives each partition of c[] from a different slave
99     if (my_rank == 0) {
100         for (int rank = 1; rank < mpi_size; ++rank) {
101             MPI_Status status;
102             MPI_Recv(c + rank * stripe * size, size * stripe, MPI_DOUBLE, rank, 0, ←
103                     MPI_COMM_WORLD,
104                     &status);
105         }
106     } else {
107         MPI_Send(c, stripe * size, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
108     }
109
110     if (my_rank == 0)
111         // stop the timer
112         cout << "elapsed time = " << timer.lap() << endl;
113
114     // results
115     if (print_option && my_rank == 0)
116         print(c, size, 'c');

```

```
116 MPI_Finalize(); // shut down MPI
117 }
```

5.2 Execution output

Check output is correct

```
1 [wyxiao_css534@cssmpi1 lab2]$ ./matrix 100 y > maxS.txt
2 [wyxiao_css534@cssmpi1 lab2]$ mpirun -n 4 ./matrix_mpi 100 y > maxM.txt
3 [wyxiao_css534@cssmpi1 lab2]$ diff maxS.txt maxM.txt
4 lc1
5 < elapsed time = 5185
6 ____
7 > elapsed time = 7133
8 [wyxiao_css534@cssmpi1 lab2]$
```

Check output the performance improvement: $774693 / 258639 = 2.9952$ times

```
1 [wyxiao_css534@cssmpi1 lab2]$ ./matrix 500
2 elapsed time = 774693
3 [wyxiao_css534@cssmpi1 lab2]$ mpirun -n 4 ./matrix_mpi 500
4 elapsed time = 258639
5 [wyxiao_css534@cssmpi1 lab2]$
```