

Program 3 Report

Yangxiao Wang

Contents

| | | |
|----------|--|----------|
| 1 | Documentation | 2 |
| 2 | Source code | 2 |
| 2.1 | InvertedIndexing.java | 2 |
| 3 | Execution output | 5 |
| 3.1 | Output analysis | 5 |
| 3.2 | Output of sequential version | 5 |
| 3.3 | Output of parallel version | 5 |
| 4 | Discussions | 6 |
| 5 | Lab Sessions 3 | 7 |
| 5.1 | Execution output | 7 |
| 6 | Reference | 7 |

1 Documentation

In this project, we are exploring the parallel technique - MapReduce. We used Inversed indexing as an example. My approach is straightforward, just map all the files with the same key (parameter) to the same reducer. And in the reducer, count the total number of the occurrence of the key. This implementation does not require combiner. When using combiner the output will have a extra 1 behind every file count (value). For example:

```
1 HDLC rfc2865.txt 22 1 rfc1122.txt 23 1
```

My opinion is that the reducer will run twice when combiner is set.

2 Source code

2.1 InvertedIndexing.java

```
1 import org.apache.hadoop.conf.*;
2 import org.apache.hadoop.fs.Path;
3 import org.apache.hadoop.io.*;
4 import org.apache.hadoop.mapred.*;
5 import org.apache.hadoop.util.*;
6
7 import java.io.IOException;
8 import java.util.*;
9
10
11 public class InvertedIndexing
12 {
13
14     public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, ↵
15         Text, Text>
16     {
17         JobConf conf;
18
19         public void configure(JobConf job)
20         {
21             this.conf = job;
22         }
23
24         public void map(LongWritable key, Text value, OutputCollector<Text, Text> output, ↵
25             Reporter reporter) throws IOException
26         {
27             // retrieve # keywords from JobConf
28             int argc = Integer.parseInt(conf.get("argc"));
29             // put args into a String array
30             Set<String> args = new HashSet();
31             // retrieve keywords
32             for (int i = 0; i < argc; i++)
33             {
34                 args.add(conf.get("keyword" + i));
35             }
36             // get the current file name
37             FileSplit fileSplit = (FileSplit) reporter.getInputSplit();
38             String filename = "" + fileSplit.getPath().getName();
39             String lines = value.toString();
40             StringTokenizer tokenizer = new StringTokenizer(lines);
41             //collect if next token match one of the args
42             while (tokenizer.hasMoreTokens())
43             {
44                 String x = tokenizer.nextToken();
45                 if (args.contains(x))
46                 {
47                     output.collect(new Text(x), new Text(filename));
```

```

48     }
49     }
50 }
51 }
52
53 public static class Reduce extends MapReduceBase implements Reducer<Text, Text, Text, ←
    Text>
54 {
55
56     public void reduce(Text key, Iterator<Text> values, OutputCollector<Text, Text> ←
        output, Reporter reporter) throws IOException
57     {
58         HashMap<String, Integer> hm = new HashMap<String, Integer>();
59         //Count the occurrence number of key in each file
60         while (values.hasNext())
61         {
62             String name = values.next().toString();
63             if (hm.containsKey(name))
64             {
65                 hm.put(name, hm.get(name) + 1);
66             }
67             else
68             {
69                 hm.put(name, 1);
70             }
71         }
72         //create Comparator to sort the result by count number
73         Comparator<java.util.Map.Entry<String, Integer>> valueComparator =
74         new Comparator<java.util.Map.Entry<String, Integer>>()
75         {
76             @Override
77             public int compare(java.util.Map.Entry<String, Integer> e1, java.util.Map.←
                Entry<String, Integer> e2)
78             {
79                 return e1.getValue() - e2.getValue();
80             }
81         };
82
83         //
84         sort the result
85         List<java.util.Map.Entry<String, Integer>> listDoc =
86         new ArrayList<java.util.Map.Entry<String, Integer>>(hm.entrySet());
87         Collections.sort(listDoc, valueComparator);
88
89         //create output string
90         StringBuilder sb = new StringBuilder();
91         for (java.util.Map.Entry<String, Integer> e : listDoc)
92         {
93             sb.append(e.getKey());
94             sb.append(" ");
95             sb.append(e.getValue());
96             sb.append(" ");
97         }
98
99         //output
100         Text docListC = new Text(sb.toString());
101         output.collect(key, docListC);
102     }
103 }
104
105 public static void main(String[] args) throws Exception
106 {
107     long time = System.currentTimeMillis();
108     JobConf conf = new JobConf(InvertedIndexing.class);
109     conf.setJobName("invertInd");
110
111     conf.setOutputKeyClass(Text.class);
112     conf.setOutputValueClass(Text.class);
113
114     conf.setMapperClass(Map.class);
115     //no need to combine because reducer already taken care of it

```

```

116 //conf.setCombinerClass(Reduce.class);
117 conf.setReducerClass(Reduce.class);
118
119 conf.setInputFormat(TextInputFormat.class);
120 conf.setOutputFormat(TextOutputFormat.class);
121
122 FileInputFormat.setInputPaths(conf, new Path(args[0]));
123 FileOutputFormat.setOutputPath(conf, new Path(args[1]));
124
125 // argc maintains #keywords
126 conf.set("argc", String.valueOf(args.length - 2));
127 for (int i = 0; i < args.length - 2; i++)
128 {
129     conf.set("keyword" + i, args[i + 2]);
130 }
131
132 JobClient.runJob(conf);
133 System.out.println("Elapsed time = " + (System.currentTimeMillis() - time) + " ms")↵
134 ;
135 }

```

3 Execution output

3.1 Output analysis

- The performance of sequential version: 253.143 s
- The performance of parallel version: 70.517 s
- Improvement: $253.143 / 70.517 = 3.5898$ times
- Cannot use diff to compare those two output because the order of the files with the same number of count is basically random. However, by comparing the result of the first two parameter HDLC and LAN and some large item with large count number, it is safe to say the outputs are the same.

3.2 Output of sequential version

```
1 HDLC rfc2865.txt 1 rfc1122.txt 1 rfc891.txt 2 rfc907.txt 2 rfc2863.txt 3 rfc1662.txt 4
2 LAN rfc2613.txt 1 rfc1044.txt 1 rfc4862.txt 1 rfc1123.txt 1 rfc2348.txt 1 rfc3461.txt 1
   rfc1661.txt 1 rfc1155.txt 2 rfc5321.txt 2 rfc2115.txt 2 rfc1629.txt 3 rfc1559.txt 3
   rfc1724.txt 3 rfc2895.txt 4 rfc1660.txt 5 rfc1213.txt 5 rfc1659.txt 5 rfc1658.txt 5
   rfc1212.txt 5 rfc1748.txt 6 rfc1694.txt 7 rfc1122.txt 10 rfc2427.txt 11 rfc950.txt 12
   rfc2067.txt 17
3 PPP rfc5531.txt 1 rfc2115.txt 1 rfc4861.txt 1 rfc1981.txt 1 rfc2427.txt 1 rfc4862.txt 1
   rfc1659.txt 1 rfc4941.txt 2 rfc5036.txt 2 rfc2460.txt 3 rfc2863.txt 12 rfc2865.txt 16
   rfc1762.txt 19 rfc1994.txt 21 rfc1662.txt 22 rfc1989.txt 26 rfc1661.txt 40 rfc5072.txt 61
   rfc1990.txt 72
4 TCP rfc6152.txt 1 rfc907.txt 1 rfc919.txt 1 rfc5322.txt 1 rfc2289.txt 1 rfc4456.txt 1
   rfc2067.txt 1 rfc922.txt 1 rfc868.txt 1 rfc5730.txt 1 rfc1155.txt 1 rfc1658.txt 1
   rfc4941.txt 1 rfc1870.txt 1 rfc3550.txt 1 rfc2355.txt 2 rfc1044.txt 2 rfc1188.txt 2
   rfc1132.txt 2 rfc1201.txt 2 rfc5065.txt 2 rfc1288.txt 2 rfc3986.txt 2 rfc1390.txt 2
   rfc894.txt 2 rfc895.txt 2 rfc1184.txt 2 rfc862.txt 3 rfc5531.txt 3 rfc863.txt 3
   rfc792.txt 3 rfc3912.txt 3 rfc3801.txt 3 rfc2895.txt 3 rfc867.txt 3 rfc1042.txt 3
   rfc866.txt 3 rfc1055.txt 3 rfc865.txt 3 rfc1356.txt 3 rfc1034.txt 5 rfc1772.txt 5
   rfc864.txt 5 rfc959.txt 5 rfc3551.txt 6 rfc4862.txt 7 rfc1939.txt 8 rfc2741.txt 8
   rfc2920.txt 9 rfc4861.txt 9 rfc854.txt 10 rfc2865.txt 10 rfc5321.txt 10 rfc2132.txt 11
   rfc791.txt 12 rfc2460.txt 12 rfc1035.txt 12 rfc1981.txt 18 rfc1006.txt 24
   rfc1191.txt 25 rfc1213.txt 33 rfc1002.txt 38 rfc1123.txt 41 rfc5734.txt 42
   rfc5036.txt 58 rfc5681.txt 83 rfc1001.txt 123 rfc4271.txt 126 rfc1122.txt 221
   rfc793.txt 278
5 UDP rfc868.txt 1 rfc1629.txt 1 rfc2348.txt 1 rfc2132.txt 1 rfc1055.txt 1 rfc950.txt 1
   rfc5531.txt 2 rfc791.txt 2 rfc4862.txt 2 rfc1034.txt 2 rfc3411.txt 2 rfc2453.txt 2
   rfc867.txt 3 rfc862.txt 3 rfc1981.txt 3 rfc1350.txt 3 rfc863.txt 3 rfc792.txt 3
   rfc1191.txt 3 rfc866.txt 3 rfc865.txt 3 rfc2895.txt 3 rfc864.txt 4 rfc3551.txt 4
   rfc4502.txt 5 rfc2131.txt 5 rfc768.txt 6 rfc2460.txt 8 rfc5036.txt 10 rfc951.txt 11
   rfc3417.txt 12 rfc1035.txt 13 rfc3550.txt 15 rfc1213.txt 19 rfc1542.txt 21 rfc2865.txt 24
   rfc1123.txt 25 rfc1001.txt 33 rfc1002.txt 50 rfc1122.txt 65
```

3.3 Output of parallel version

```
1 HDLC rfc2865.txt 1 rfc1122.txt 1 rfc891.txt 2 rfc907.txt 2 rfc2863.txt 3 rfc1662.txt 4
2 LAN rfc2613.txt 1 rfc1044.txt 1 rfc4862.txt 1 rfc1123.txt 1 rfc2348.txt 1 rfc3461.txt 1
   rfc1661.txt 1 rfc1155.txt 2 rfc5321.txt 2 rfc2115.txt 2 rfc1629.txt 3 rfc1559.txt 3
   rfc1724.txt 3 rfc2895.txt 4 rfc1660.txt 5 rfc1213.txt 5 rfc1659.txt 5 rfc1658.txt 5
   rfc1212.txt 5 rfc1748.txt 6 rfc1694.txt 7 rfc1122.txt 10 rfc2427.txt 11 rfc950.txt 12
   rfc2067.txt 17
3 PPP rfc5531.txt 1 rfc2115.txt 1 rfc4861.txt 1 rfc1981.txt 1 rfc2427.txt 1 rfc4862.txt 1
   rfc1659.txt 1 rfc4941.txt 2 rfc5036.txt 2 rfc2460.txt 3 rfc2863.txt 12 rfc2865.txt 16
   rfc1762.txt 19 rfc1994.txt 21 rfc1662.txt 22 rfc1989.txt 26 rfc1661.txt 40 rfc5072.txt 61
   rfc1990.txt 72
4 TCP rfc6152.txt 1 rfc907.txt 1 rfc919.txt 1 rfc5322.txt 1 rfc2289.txt 1 rfc4456.txt 1
   rfc2067.txt 1 rfc922.txt 1 rfc868.txt 1 rfc5730.txt 1 rfc1155.txt 1 rfc1658.txt 1
   rfc4941.txt 1 rfc1870.txt 1 rfc3550.txt 1 rfc2355.txt 2 rfc1044.txt 2 rfc1188.txt 2
```

```

    rfc1132.txt 2 rfc1201.txt 2 rfc5065.txt 2 rfc1288.txt 2 rfc3986.txt 2 rfc1390.txt ←
2 rfc894.txt 2 rfc895.txt 2 rfc1184.txt 2 rfc862.txt 3 rfc5531.txt 3 rfc792.txt 3 ←
rfc863.txt 3 rfc3912.txt 3 rfc3801.txt 3 rfc2895.txt 3 rfc867.txt 3 rfc1042.txt 3 ←
rfc866.txt 3 rfc1055.txt 3 rfc865.txt 3 rfc1356.txt 3 rfc1034.txt 5 rfc1772.txt 5 ←
rfc959.txt 5 rfc864.txt 5 rfc3551.txt 6 rfc4862.txt 7 rfc1939.txt 8 rfc2741.txt 8 ←
rfc2920.txt 9 rfc4861.txt 9 rfc854.txt 10 rfc2865.txt 10 rfc5321.txt 10 rfc2132.txt ←
11 rfc791.txt 12 rfc2460.txt 12 rfc1035.txt 12 rfc1981.txt 18 rfc1006.txt 24 ←
rfc1191.txt 25 rfc1213.txt 33 rfc1002.txt 38 rfc1123.txt 41 rfc5734.txt 42 rfc5036 ←
txt 58 rfc5681.txt 83 rfc1001.txt 123 rfc4271.txt 126 rfc1122.txt 221 rfc793.txt ←
278
5 UDP rfc868.txt 1 rfc1629.txt 1 rfc2348.txt 1 rfc2132.txt 1 rfc1055.txt 1 rfc950.txt 1 ←
rfc5531.txt 2 rfc791.txt 2 rfc4862.txt 2 rfc1034.txt 2 rfc3411.txt 2 rfc2453.txt 2 ←
rfc862.txt 3 rfc867.txt 3 rfc1981.txt 3 rfc1350.txt 3 rfc863.txt 3 rfc792.txt 3 ←
rfc1191.txt 3 rfc866.txt 3 rfc865.txt 3 rfc2895.txt 3 rfc864.txt 4 rfc3551.txt 4 ←
rfc4502.txt 5 rfc2131.txt 5 rfc768.txt 6 rfc2460.txt 8 rfc5036.txt 10 rfc951.txt 11 ←
rfc3417.txt 12 rfc1035.txt 13 rfc3550.txt 15 rfc1213.txt 19 rfc1542.txt 21 rfc2865 ←
.txt 24 rfc1123.txt 25 rfc1001.txt 33 rfc1002.txt 50 rfc1122.txt 65

```

4 Discussions

1. File distribution over a cluster system

MPI is Message Passing Interface, it does not need a file system to store its data. And the data is sent to another node to be computed.

On the other hand, MapReduce is usually used with Hadoop Distributed File System. And the data is stored in local storage on each data node.

2. Collective/Reductive operation to create inverted indexing

The implementation with MPI will be:

- read all data, distribute the data to each node.
- each node will compute and find the number of keywords' occurrence in their portion of data.
- share the whole results and do the reduction.

The biggest problem with MPI when doing inverted indexing is that it uses network to transfer data. When the size of files is large, the MPI's performance will not be ideal.

3. Amount of boilerplate code

Comparing to MapReduce, MPI would have more boilerplate code like send, receive, and parse data. And MapReduce only need implement the Map class and Reducer class.

4. Anticipated execution performance

Again, the limitation of MPI is the network, when the size of data/file is large and the performance could be slower than MapReduce. However, if the data is not too large, using MPI could be more efficient because the Hadoop's fault tolerance system can slow down the process.

5. Fault tolerance; recovery from a crash

MPI support checkpoint to restart from the checkpoint if anything goes wrong. It does not have Message logging techniques, data Reliability and network fault tolerance, User directed and communicator driven fault tolerance. Basically, developer can set multiple checkpoints before passing data or during iteration. However, the data could be still lost if the network is not stable.

Hadoop has its own built-in fault tolerance and fault compensation capabilities. Every data block has a copy that is stored on other servers. And it also generate logs during the execution process.

5 Lab Sessions 3

5.1 Execution output

```
[wyxiao_css534@cssmpi1 wordcount_2.0]$ hadoop jar wordcount.jar WordCount input output
18/11/08 18:37:12 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should implement Tool for the same.
18/11/08 18:37:12 INFO mapred.FileInputFormat: Total input paths to process : 2
18/11/08 18:37:13 INFO mapred.JobClient: Running job: job_201811071847_0013
18/11/08 18:37:14 INFO mapred.JobClient: map 0% reduce 0%
18/11/08 18:37:21 INFO mapred.JobClient: map 33% reduce 0%
18/11/08 18:37:23 INFO mapred.JobClient: map 66% reduce 0%
18/11/08 18:37:28 INFO mapred.JobClient: map 100% reduce 0%
18/11/08 18:37:30 INFO mapred.JobClient: map 100% reduce 11%
18/11/08 18:37:33 INFO mapred.JobClient: map 100% reduce 22%
18/11/08 18:37:39 INFO mapred.JobClient: map 100% reduce 100%
18/11/08 18:37:41 INFO mapred.JobClient: Job complete: job_201811071847_0013
18/11/08 18:37:41 INFO mapred.JobClient: Counters: 18
18/11/08 18:37:41 INFO mapred.JobClient:   Map-Reduce Framework
18/11/08 18:37:41 INFO mapred.JobClient:     Combine output records=6
18/11/08 18:37:41 INFO mapred.JobClient:     Spilled Records=12
18/11/08 18:37:41 INFO mapred.JobClient:     Reduce input records=6
18/11/08 18:37:41 INFO mapred.JobClient:     Reduce output records=5
18/11/08 18:37:41 INFO mapred.JobClient:     Map input records=2
18/11/08 18:37:41 INFO mapred.JobClient:     Map output records=8
18/11/08 18:37:41 INFO mapred.JobClient:     Map output bytes=82
18/11/08 18:37:41 INFO mapred.JobClient:     Reduce shuffle bytes=46
18/11/08 18:37:41 INFO mapred.JobClient:     Combine input records=8
18/11/08 18:37:41 INFO mapred.JobClient:     Map input bytes=50
18/11/08 18:37:41 INFO mapred.JobClient:     Reduce input groups=5
18/11/08 18:37:41 INFO mapred.JobClient:   FileSystemCounters
18/11/08 18:37:41 INFO mapred.JobClient:     HDFS_BYTES_READ=54
18/11/08 18:37:41 INFO mapred.JobClient:     FILE_BYTES_WRITTEN=266
18/11/08 18:37:41 INFO mapred.JobClient:     FILE_BYTES_READ=79
18/11/08 18:37:41 INFO mapred.JobClient:     HDFS_BYTES_WRITTEN=41
18/11/08 18:37:41 INFO mapred.JobClient:   Job Counters
18/11/08 18:37:41 INFO mapred.JobClient:     Launched map tasks=3
18/11/08 18:37:41 INFO mapred.JobClient:     Launched reduce tasks=1
18/11/08 18:37:41 INFO mapred.JobClient:     Rack-local map tasks=3
[wyxiao_css534@cssmpi1 wordcount_2.0]$
```

Figure 1: MapReduce execution

```
[wyxiao_css534@cssmpi1 ~]$ hadoop fs -ls /user/wyxiao_css534/output
Found 2 items
drwxr-xr-x  - wyxiao_css534 wyxiao_css534      0 2018-11-07 21:37 /user/wyxiao_css534/output/_logs
-rw-r--r--  3 wyxiao_css534 wyxiao_css534    2283 2018-11-07 21:38 /user/wyxiao_css534/output/part-00000
```

Figure 2: /user/yourAccount/output

```
[wyxiao_css534@cssmpi1 output]$ cat part-00000
Bye      1
Goodbye  1
Hadoop   2
He        1
Hello    2
World    2
he        2
```

Figure 3: part-00000